

RL Assignment 1 Programming Question

Question: Modified Gradient Bandit with Adaptive Baseline

Problem Statement

Implement a **Gradient Bandit Algorithm** with an **adaptive baseline** that automatically adjusts based on the variance of recent rewards. Your implementation should include the following components:

Background: The standard gradient bandit algorithm uses a baseline (typically the average reward) to reduce variance in updates. However, when reward distributions have different variances across arms or change over time, a simple average may not be optimal.

Task

Implement a gradient bandit agent with the following specifications:

1. **Standard Gradient Bandit Update:**
 - a. For selected action: $H[a] = H[a] + \alpha(R - \text{baseline})(1 - \pi[a])$
 - b. For non-selected actions: $H[a] = H[a] - \alpha(R - \text{baseline})\pi[a]$
 - c. Action probabilities: $\pi[a] = \exp(H[a]) / \sum \exp(H[b])$
2. **Adaptive Baseline:** Instead of using just the average reward, use a **weighted baseline** that adapts based on recent reward variance:

$$\text{baseline_t} = (1 - \beta) * \text{avg_reward} + \beta * \text{recent_variance_adjusted_mean}$$

where:

- *avg_reward* is the running average of all rewards
- *recent_variance_adjusted_mean* considers the last W rewards with weights inversely proportional to their squared deviation from the mean
- β is an adaptation parameter ($0 \leq \beta \leq 1$)

3. Performance Tracking:

- a. Track the percentage of optimal actions selected
- b. Track cumulative regret (difference from optimal arm)
- c. Compare against standard gradient bandit ($\beta = 0$)

Requirements

1. Implement the *AdaptiveGradientBandit* class with methods:
 - a. `__init__(n_arms, alpha, beta, window_size)`
 - b. `select_action()` - returns action based on softmax policy
 - c. `update(action, reward)` - updates preferences using adaptive baseline
 - d. `get_baseline()` - returns current baseline value
2. Run experiments on a **non-stationary 10-armed bandit** where:
 - a. True reward means start at random values from $N(0, 1)$
 - b. Every 500 steps, add $N(0, 0.1)$ to each arm's true mean (random walk)
 - c. Actual rewards are sampled from $N(\text{true_mean}, 1)$
3. Compare performance over 2000 steps for:
 - a. Standard gradient bandit ($\beta = 0$)
 - b. Adaptive gradient bandit ($\beta = 0.3$)
 - c. Adaptive gradient bandit ($\beta = 0.6$)
4. Plot:
 - a. Average reward over time (running average over last 100 steps)
 - b. Percentage of optimal action selection
 - c. Baseline values over time for different β values

Function Signatures:

```
// Add function signatures here.
```

Expected Observations

- Adaptive baselines should perform better in non-stationary environments
- Higher β values should show faster adaptation but potentially more variance
- Baseline values should track the changing reward landscape